

✓ Support Vector Machine (SVM) from Scratch Using Hinge Loss and Subgradient

This notebook demonstrates the implementation of a linear Support Vector Machine (SVM) classifier **from scratch**, without using scikit-learn's built-in SVM. The model is trained using **hinge loss** and both **gradient descent (GD)** and **subgradient descent** optimization techniques.

Hinge Loss

The hinge loss is commonly used for training "maximum-margin" classifiers, particularly SVMs. It is defined as:

The hinge loss $\mathcal{L}_{\text{hinge}}$ for a single example (x_i, y_i) with feature vector x_i and true label $y_i \in \{-1, +1\}$ is given by:

$$\mathcal{L}_{\text{hinge}}(x_i, y_i; \mathbf{w}, b) = \max(0, 1 - y_i(\mathbf{w}^T x_i + b))$$

where:


- \mathbf{w} is the weight vector (coefficients) of the hyperplane,
- b is the bias term,
- y_i is the true label (+1 or -1),
- $\mathbf{w}^T x_i + b$ is the decision function output.

SVM Objective Function

SVM aims to minimize the regularized hinge loss:

Visualization

SVM: combing back to hinge loss



Minimize: $g(\mathbf{w}) = \frac{1}{2} \|\mathbf{w}\|^2 + C \times \sum_{i=1}^n (1 - y_i \times (\mathbf{w} \cdot \mathbf{x}_i + b))_+$

$$\frac{\partial g(\mathbf{w})}{\partial \mathbf{w}} = \mathbf{w} + C \times \sum_{i=1}^n \begin{cases} 0 & \text{if } y_i \times (\mathbf{w} \cdot \mathbf{x}_i + b) \geq 1 \\ -y_i \times \mathbf{x}_i & \text{otherwise} \end{cases}$$

27

Here is a visualization of the hinge loss function and the effect of the regularization term in the SVM objective.

- ✓ Imports essential libraries for data handling, visualization, evaluation, and splitting. `svc` is imported for comparison at the end.

Double-click (or enter) to edit

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np
from sklearn.metrics import accuracy_score, confusion_matrix
sns.set_theme()
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
```

Imports essential libraries for data handling, visualization,

- ✓ evaluation, and splitting. `svc` is imported for comparison at the end.

```
def wrangle(path):
    data=pd.read_csv(path)
    data.isnull().sum()
    data.dropna(inplace=True)
    data.duplicated().sum()
    data.drop(columns='rad',inplace=True)
    data.rename(columns={'chas': 'target'}, inplace=True)
    data['target'] = data['target'].replace(0, -1)
    return data
```

- ✓ Loads the dataset using the `wrangle` function and displays the cleaned DataFrame.

```
data=wrangle(r"BostonHousing.csv")
data
```



	crim	zn	indus	target	nox	rm	age	dis	tax	ptratio	b	lsta
0	0.00632	18.0	2.31	-1	0.538	6.575	65.2	4.0900	296	15.3	396.90	4.9
1	0.02731	0.0	7.07	-1	0.469	6.421	78.9	4.9671	242	17.8	396.90	9.1
2	0.02729	0.0	7.07	-1	0.469	7.185	61.1	4.9671	242	17.8	392.83	4.0
3	0.03237	0.0	2.18	-1	0.458	6.998	45.8	6.0622	222	18.7	394.63	2.9
4	0.06905	0.0	2.18	-1	0.458	7.147	54.2	6.0622	222	18.7	396.90	5.3
...
501	0.06263	0.0	11.93	-1	0.573	6.593	69.1	2.4786	273	21.0	391.99	9.6
502	0.04527	0.0	11.93	-1	0.573	6.120	76.7	2.2875	273	21.0	396.90	9.0
503	0.06076	0.0	11.93	-1	0.573	6.976	91.0	2.1675	273	21.0	396.90	5.6
504	0.10959	0.0	11.93	-1	0.573	6.794	89.3	2.3889	273	21.0	393.45	6.4
505	0.04741	0.0	11.93	-1	0.573	6.030	80.8	2.5050	273	21.0	396.90	7.8

501 rows × 13 columns

Generates a correlation matrix heatmap to visualize feature


- ✓ relationships and help identify potentially redundant or

correlated features.

```
plt.figure(figsize=(12,8))
sns.heatmap(data.corr(),annot=True,cmap='coolwarm')
```



<Axes: >




- ✓ Displays the class distribution in the target variable to understand class balance.

```
data['target'].value_counts(normalize=True).plot(kind='bar')
```





<Axes: xlabel='target'>



- ✓ Plots a scatter plot between `age` and `crim` to visually inspect any potential pattern or correlation.

```
plt.scatter(data['age'],data['crim'])
```

 <matplotlib.collections.PathCollection at 0x7c602d186450>



- Separates features (X) and target (y), converts them to NumPy arrays, and
- ✓ performs an 80-20 train-test split. Further splits training data into training and validation sets.

```
X=data.drop(columns='target')
y=data['target']
X=X.to_numpy()
y=y.to_numpy()
x_train_val,x_test,y_train_val,y_test=train_test_split(X,y,test_size=0.2)
x_train,x_val,y_train,y_val=train_test_split(x_train_val,y_train_val,test_size=0.2)
```

- ✓ Defines a custom SVM class from scratch with the following methods:

- `__init__`: Sets hyperparameters (learning rate, lambda, epochs, method).
- `loss_calc`: Calculates hinge loss with L2 regularization.
- `accuracy`: Calculates classification accuracy.
- `fit`: Trains the SVM using either gradient descent (`gd`) or subgradient (`subgd`).
- `predict`: Predicts using the learned weights and bias.

```
class svm_scratch:
    def __init__(self,lr=0.01,lambda_param=0.01,epochs=100,method='gd'):
```

```

self.lr=lr
self.method=method
self.lambda_param=lambda_param
self.epochs=epochs
self.w=None
self.b=None
self.loss_history=[]
self.acc_train=[]
self.val_train=[]

def loss_calc(self,X,y):
    margin=y*(np.dot(X,self.w)+self.b)
    loss=0.5*np.dot(self.w,self.w)+np.mean(np.maximum(0,1-margin))
    return loss

def accuracy(self,X,y):
    predict=np.sign(np.dot(X,self.w)+self.b)
    return accuracy_score(y_true=y,y_pred=predict)

def fit(self,X_train,y_train,X_val=None,y_val=None):
    n_sample,n_feature=X_train.shape
    self.w=np.zeros(n_feature)
    self.b=0

    for i in range(0,self.epochs):
        for idx,x_i in enumerate(X_train):
            condition=y_train[idx]*(np.dot(x_i ,self.w)+ self.b)

            if self.method == 'gd':
                if condition >=1 :
                    dw= 2*self.lambda_param*self.w
                    db=0
                elif condition < 1:
                    dw=2*self.lambda_param*self.w-np.dot(x_i,y_train[idx])
                    db=-y_train[idx]
            elif self.method=='subgd':
                if condition >1 :
                    dw= 2*self.lambda_param*self.w
                    db=0
                elif condition < 1:
                    dw=2*self.lambda_param*self.w-np.dot(x_i,y_train[idx])
                    db=-y_train[idx]
                elif condition==1:
                    rand_int=np.random.uniform(0,1)
                    dw=2*self.lambda_param *self.w + rand_int*(-y_train[idx] * x_i)
                    db=rand_int *-y_train[idx]
            self.w=self.w-self.lr*dw
            self.b=self.b-self.lr*db
        print(f'loss at epoch {i} = {self.loss_calc(X_train,y_train)}')

        self.loss_history.append(self.loss_calc(X_train,y_train))

        self.acc_train.append(self.accuracy(X_train, y_train))
        if X_val is not None and y_val is not None:
            self.val_train.append(self.accuracy(X_val, y_val))

```

```
def predict(self,X):
    output=np.dot(X,self.w)+self.b
    return np.sign(output)
```

✓ Initializes and trains two SVM models:

- One using **subgradient descent**
- Another using **gradient descent** (with validation monitoring).

```
svm_subgrad = svm_scratch(lr=0.001, lambda_param=0.01, epochs=100, method="subgd")
svm_subgrad.fit(x_train, y_train)
```

```
svm_grad = svm_scratch(lr=0.001, lambda_param=0.01, epochs=100, method="gd")
svm_grad.fit(x_train, y_train, x_val, y_val)
```

```
→ loss at epoch 0 = 27.256915753737882
loss at epoch 1 = 23.112916639626143
loss at epoch 2 = 16.00292610716297
loss at epoch 3 = 14.739522392692805
loss at epoch 4 = 23.505143132624536
loss at epoch 5 = 22.800162264828938
loss at epoch 6 = 26.766030101463933
loss at epoch 7 = 25.66488301919188
loss at epoch 8 = 27.738123958424605
loss at epoch 9 = 26.64322641767932
loss at epoch 10 = 27.83496832546606
loss at epoch 11 = 26.77556381729994
loss at epoch 12 = 27.99029950001002
loss at epoch 13 = 27.96288849235991
loss at epoch 14 = 26.930331530878654
loss at epoch 15 = 29.26159569923629
loss at epoch 16 = 27.758587018848424
loss at epoch 17 = 27.052053564066554
loss at epoch 18 = 27.975915777443685
loss at epoch 19 = 28.9001817144114
loss at epoch 20 = 28.460513120001565
loss at epoch 21 = 28.64342256584927
loss at epoch 22 = 28.996777408676678
loss at epoch 23 = 26.806751428843423
loss at epoch 24 = 22.56147110647156
loss at epoch 25 = 24.149753313123966
loss at epoch 26 = 32.21860341743858
loss at epoch 27 = 32.605253563870974
loss at epoch 28 = 30.35402297325654
loss at epoch 29 = 30.722235226897524
loss at epoch 30 = 24.700211115206564
loss at epoch 31 = 14.297079493299535
loss at epoch 32 = 32.92126808345259
```



```
loss at epoch 33 = 33.36249785980743
loss at epoch 34 = 13.978914017249028
loss at epoch 35 = 31.753500806135595
loss at epoch 36 = 27.30150683044185
loss at epoch 37 = 30.99840725142428
loss at epoch 38 = 32.02880763707009
loss at epoch 39 = 30.71501133349245
loss at epoch 40 = 26.904006894026075
loss at epoch 41 = 31.63125030753334
loss at epoch 42 = 26.10546044856649
loss at epoch 43 = 34.20662793512189
loss at epoch 44 = 25.726565942419327
loss at epoch 45 = 28.815239072376947
loss at epoch 46 = 34.33593019880354
loss at epoch 47 = 26.53701755198845
loss at epoch 48 = 35.36873343352303
loss at epoch 49 = 28.935106479138845
loss at epoch 50 = 27.33039957481248
loss at epoch 51 = 28.85015836503139
loss at epoch 52 = 28.033516065634327
loss at epoch 53 = 16.1133110230466
loss at epoch 54 = 35.355265793562765
loss at epoch 55 = 27.48451379621134
loss at epoch 56 = 16.487542401737016
loss at epoch 57 = 33.66457828417543
```

- Similar to the previous cell but evaluates the gradient descent-
✓ based SVM model. Shows final accuracies and plots the loss curve.

```
# Loss per epoch
losses_sub = svm_subgrad.loss_history
```

```
# Training accuracy per epoch
train_accuracies_sub = svm_subgrad.acc_train
```


```
# Validation accuracy per epoch
val_accuracies_sub = svm_subgrad.val_train
```

```
print("the Training Accuracy:", train_accuracies_sub[-1]*100)
y_pred_sub=svm_subgrad.predict(x_val)
print(f'the Validation Accuracy = {accuracy_score(y_val,y_pred_sub)*100}')
y_pred_sub=svm_subgrad.predict(x_test)
print(f'the Test Accuracy = {accuracy_score(y_test,y_pred_sub)*100}')
```

```
plt.plot(losses_sub)
plt.title("Loss Over Epochs")
plt.xlabel("Epoch")
```

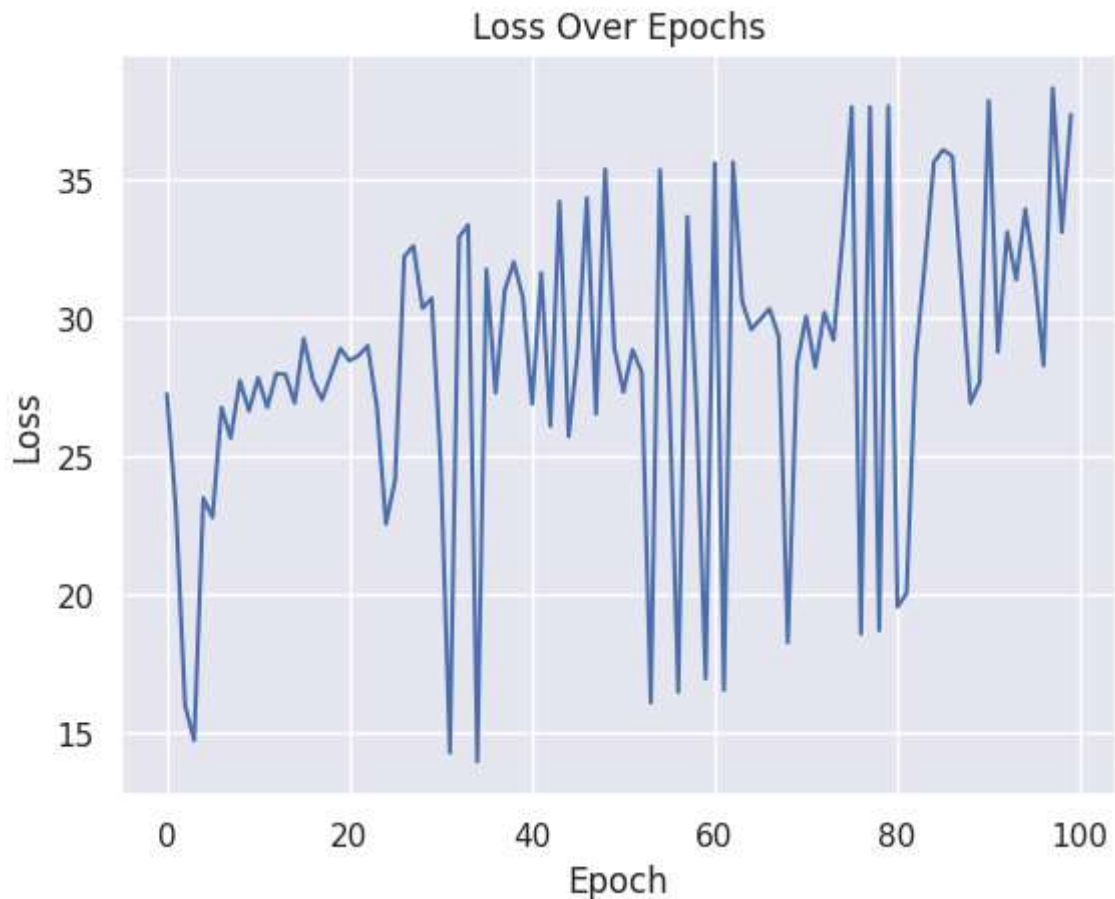
```
plt.ylabel("Loss")
plt.show()
```

⇒ the Training Accuracy: 92.5
 the Validation Accuracy = 96.25
 the Test Accuracy = 92.07920792079209



```
# Loss per epoch
losses_grad = svm_grad.loss_history
# Training accuracy per epoch
train_accuracies_grad = svm_grad.acc_train
# Validation accuracy per epoch
val_accuracies_grad = svm_grad.val_train
print("the Training Accuracy:", train_accuracies_grad[-1]*100)
y_pred_grad=svm_grad.predict(x_val)
print(f'the Validation Accuracy = {accuracy_score(y_val,y_pred_grad)*100}')
y_pred_grad=svm_grad.predict(x_test)
print(f'the Test Accuracy = {accuracy_score(y_test,y_pred_grad)*100}')
plt.plot(losses_grad)
plt.title("Loss Over Epochs")
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.show()
```

the Training Accuracy: 92.5
 the Validation Accuracy = 96.25
 the Test Accuracy = 92.07920792079209



Fits a standard SVM model from scikit-learn using a linear

- ✓ kernel and compares its test accuracy with the custom implementation.

```

svc_model = SVC(kernel='linear', C=1.0)
svc_model.fit(x_train, y_train)
y_pred_sklearn = svc_model.predict(x_test)
print(f'the Test Accuracy = {accuracy_score(y_test, y_pred_sklearn)*100}')

```

the Test Accuracy = 92.07920792079209

Displays a confusion matrix heatmap for the gradient descent-based model on the test set to visualize classification performance.

✓ For confusion_matrix

```
# === Generate predictions ===
y_pred_sub = svm_subgrad.predict(x_test)
y_pred_grad = svm_grad.predict(x_test)
y_pred_sklearn = svc_model.predict(x_test)

# === Compute confusion matrices ===
conf_matrix_sub = confusion_matrix(y_test, y_pred_sub)
conf_matrix_grad = confusion_matrix(y_test, y_pred_grad)
conf_matrix_svc = confusion_matrix(y_test, y_pred_sklearn)


# === Create subplots ===
fig, axs = plt.subplots(1, 3, figsize=(18, 5))

# Plot Subgradient Descent SVM
sns.heatmap(conf_matrix_sub, annot=True, fmt='d', cmap='coolwarm',
            xticklabels=['-1', '1'], yticklabels=['-1', '1'], ax=axs[0])
axs[0].set_title('Subgradient Descent SVM')
axs[0].set_xlabel('Predicted')
axs[0].set_ylabel('Actual')

# Plot Gradient Descent SVM
sns.heatmap(conf_matrix_grad, annot=True, fmt='d', cmap='coolwarm',
            xticklabels=['-1', '1'], yticklabels=['-1', '1'], ax=axs[1])
axs[1].set_title('Gradient Descent SVM')
axs[1].set_xlabel('Predicted')
axs[1].set_ylabel('Actual')

# Plot Sklearn SVC
sns.heatmap(conf_matrix_svc, annot=True, fmt='d', cmap='coolwarm',
            xticklabels=['-1', '1'], yticklabels=['-1', '1'], ax=axs[2])
axs[2].set_title('SVC Model')
axs[2].set_xlabel('Predicted')
axs[2].set_ylabel('Actual')

plt.tight_layout()
plt.show()
```



✓ Compare Final Accuracy Across Models

```
test_accuracies = [
    accuracy_score(y_test, y_pred_sub),
    accuracy_score(y_test, y_pred_grad),
    accuracy_score(y_test, y_pred_sklern)
]
labels = ['SubGD SVM', 'GD SVM', 'SVC']
accuracy_percentages = [acc * 100 for acc in test_accuracies]

plt.figure(figsize=(8,6))
sns.barplot(x=labels, y=accuracy_percentages, palette="viridis")
plt.title("Test Accuracy Comparison")
plt.ylabel("Accuracy (%)")
plt.ylim(0, 100)
plt.grid(axis='y')

# Add value labels on bars
for i, acc in enumerate(accuracy_percentages):
    plt.text(i, acc + 1, f"{acc:.2f}%", ha='center', va='bottom', fontweight='bold')

plt.show()
```



<ipython-input-18-0342ea392552>:10: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.

```
sns.barplot(x=labels, y=accuracy_percentages, palette="viridis")
```